

Computable Compressed Matrices

Crysttian Arantes Paixão^{1*}, Flávio Codeço Coelho²,

1 Crysttian Arantes Paixão Applied Mathematics School, Getulio Vargas Foundation, Rio de Janeiro, RJ, Brazil

2 Flávio Codeço Coelho Applied Mathematics School, Getulio Vargas Foundation, Rio de Janeiro, RJ, Brazil

*** E-mail: Corresponding author@institute.edu**

Abstract

Author Summary

Introduction

Data compression is traditionally used to reduce storage resources usage and/or transmission costs [1]. Compression techniques can be classified into lossy and lossless. Examples of lossy data compression are MP3 (audio), JPEG (image) and MPEG (video). In this paper we discuss the use of lossless compression for numerical data structures such as numerical arrays to achieve compression without losing the mathematical properties of the original data.

Lossless compression methods usually exploit redundancies present in the data in order to find a shorter form of describing the same information content. For example, a dictionary-based compression, only stores the positions in which a given word occurs in a document, thus saving the space required to store all its repetitions [2].

Any kind of compression incurs some computational cost. Such costs often have to be paid twice since the data needs to be decompressed to be used for its original purpose. Sometimes computational costs are irrelevant, but the need to decompress for usage, can signify that the space saved with compression must be available when data is decompressed for usage, thus partially negating the advantages of compression.

Most if not all existing lossless compression methods were developed under the following usage paradigm: *produce* \rightarrow *compress* \rightarrow *store* \rightarrow *uncompress* \rightarrow *use*. The focus of the present work is to allow a slightly different usage: *produce* \rightarrow *compress* \rightarrow *perform mathematical manipulations to use* and *decompress* only for human reading.

With the growth of data volumes and analytical demands, creative solutions are needed to efficiently store as well as consume it on demand. This issue is present in many areas of application, ranging from business to science [3], and is being called the *Big Data* phenomenon. In the world of Big Data, the need to analyze data immediately after its coming into existence became the norm. And this analysis must take place, efficiently, within the confines of (RAM) memory. This kind of analyses are what is now known as streaming data analysis [4]. Given a sufficiently dense stream of data, compression and decompression costs may become prohibitive. So having a way to compress data and keeping it compressed for the entire course of the analytical pipeline, is very desirable.

This paper will focus solely on numerical data which for the purpose of the applications is organized as matrices. This is a most common data structure found in computational data analysis environments. The matrices compressed according to the methodologies proposed here should be able to undergo the same mathematical operations as the original uncompressed matrices. e.g. linear algebra manipulations. This way, the cost of compression is reduced to a single event of compression and no need of decompression except when displaying the results for human reading. The idea of operating with compressed arrays is relatively new [5], and it has yet to find mainstream applications to the field of numerical computations. One application which employs a form of compression is the sparse matrix linear algebra algorithms [6], in this case there is no alteration in the standard encoding of the data, but only the non-zero elements

of the matrices are stored and operated upon.

Larger than RAM data structures can render traditional analytical algorithms impracticable. Currently, the technique most commonly used when dealing with large matrices for numerical computations, is memory mapping [7,8]. In memory mapping the matrix is allocated in a virtual contiguous address space which extends from memory into disk. Thus, larger than memory data structures can be manipulated as if they were in memory. This technique has a big performance penalty due to lower access speeds of disk when compared to RAM.

In this paper we present two methods for the lossless compression of (numerical) arrays. The methods involve the encoding of the numbers as strings of bits of variable length. The methods resemble the arithmetic coding [9] algorithm, but is cheaper to compute. We describe the process of compression and decompression, and study their efficiency under different applications. We also discuss the efficiency of the compression as a function of the distribution of the elements of the matrix.

Methods

Matrix compression

To maintain mathematical equivalence with the original data for any arithmetic operations, we need to maintain the structure of the matrix, i.e., the ability to access any element given its row i and column j and also the numeric nature of its elements. In order to achieve compression we decided to exploit inefficiencies in the conventional way matrices are allocated in memory. The examples in this paper will be restricted to matrices with positive integer elements.

The compression method is as follows. Let $M_{r \times c}$ be a matrix, in which r is the number of rows and c the number of columns. Each element of this matrix, called m_{ij} , is a positive integer. In digital computers, all information is stored as binary code (base 2 numbers). However the conventional way to store arrays of integers is on a memory block sequence of fixed size (power of 2 bit numbers), one for each element. The maximum size of a block is equal to the word size of the processor, which for most current CPUs is 64 bits. Some special number such as complex number may be encoded as two blocks instead of one. The size of the chunk of memory allocated to each number will determine their maximum size (for integers) or their precision (for floating-point numbers). So for matrix M , the total memory allocated, assuming chunks of 64 bits, is given by $\mathcal{B} = r \times c \times 64$.

The number of bits allocated \mathcal{B} , is larger than the absolute minimum number of bits required to represent all the elements of M , since smaller integers, when converted to base 2, require less digits. From now on, when the numerical base will be explicitly notated when necessary to avoid confusion between binary and decimal integers.

Let's consider an extreme example: a matrix composed exclusively of 0s and 1s (base 10). If the matrix type is set to 64-bit integers, 63 bits will be wasted per element of the matrix, since the minimum number of bits needed to store such a matrix is $\mathfrak{b} = r \times c \times 1$. The potential economy of bits ξ can be represented by $\xi = \mathcal{B} - \mathfrak{b} = r \times c \times 63$.

So it is evident that for any matrix whose greatest element requires less than 64 bits (or the fixed type of the matrix) to be represented, potential memory savings will grow linearly with the size of the matrix.

Method 1: The Supreme Minimum (SM)

The SM method consists in determining the value of the greatest element of matrix M , which coincides with its supremum, $\max M == \sup M$ and determine the minimum bit number, $b(\sup M)$ (Equation 1), required to store it. We will use capital roman letters to denote uncompressed matrices and the corresponding lower case letter for the compressed version.

As we see in 5, η_1 does not depend on size of the matrix, only on the bit-length of $\max M$. If $b(\max M) = 64$, η_1 is 0, i. e., no compression is possible. On the other extreme, if the matrix is composed exclusively of 0s and 1s, maximal compression is achievable, $\eta_1 = 1$.

VLB Method

For the VLB method, compression depends on the value of each element of the matrix. In this method bit-length variability affects the compression ratio, so the formula will have to include this information.

Let the rc elements of the matrix $M_{r \times c}$ be divided into g groups, each with f_i numbers of bit-length $b_i = b(m_i)$. Thus f_i is the frequency of each bit-length present in M . Let $k = b(b(\max M))$, i.e., the bit length of the bit-length of $\max M$. The efficiency η_2 is shown below.

$$\eta_2 = \frac{64 \times rc - \sum_{i=1}^g (b_i + k) \times f_i}{64 \times rc} \quad (6)$$

We can further simplify Equation 6 to get a shorter expression for the compression ratio.

$$\begin{aligned} \eta_2 &= \frac{64 \times rc - \sum_{i=1}^g (b_i \times f_i + k \times f_i)}{64 \times rc} \\ &= \frac{64 \times rc - \sum_{i=1}^g b_i \times f_i - \sum_{i=1}^g k \times f_i}{64 \times rc} \\ &= \frac{64 \times rc - \sum_{i=1}^g b_i \times f_i - k \times \sum_{i=1}^g f_i}{64 \times rc} \end{aligned}$$

Please note that $\sum_{i=1}^g f_i = rc$,

$$\begin{aligned} \eta_2 &= \frac{64 \times rc - \sum_{i=1}^g b_i \times f_i - k \times rc}{64 \times rc} \\ &= 1 - \frac{\sum_{i=1}^g b_i \times f_i}{64 \times rc} - \frac{k}{64} \end{aligned} \quad (7)$$

Results

Random Matrix Generation. In both methods, compression efficiency depends on the distribution of the bit-lengths $b(m_{i,j})$. Thus, in this section, a method to generate a variety of random bit-length distributions is proposed.

For simplicity we will model the distribution of as a mixture X of two Beta distributions, $B_1 \sim \text{Beta}(\alpha_1, \beta_1)$ and $B_2 \sim \text{Beta}(\alpha_2, \beta_2)$, whose probability function is shown in Equation 8. Since the Beta distribution is defined only in the interval $[0, 1] \subset \mathbb{R}$, we applied a simple transformation $(\lfloor 64 \times x \rfloor + 1)$ to the mixture in order to map it to the interval of $[1, 64] \subset \mathbb{Z}$.

$$F(x) = w \text{Beta}(\alpha_1, \beta_1) + (1 - w) \text{Beta}(\alpha_2, \beta_2) \quad (8)$$

The intention of using this mixture was to find a simple way to represent a large variety of bit length distributions. The first two central moments of this mixture are given in 9 and will be used later to summarize our numerical results.

$$\begin{aligned}
E(X) &= wE(B_1) + (1-w)E(B_2) \\
&= w \frac{\alpha_1}{\alpha_1 + \beta_1} + (1-w) \frac{\alpha_2}{\alpha_2 + \beta_2} \\
Var(X) &= wVar(B_1) + (1-w)Var(B_2) + w(1-w)(E(B_1)^2 - E(B_2)^2)
\end{aligned} \tag{9}$$

In order to explore the compression efficiency of both methods, we generated samples from the mixture defined above, varying its parameters. From now on, when we mention Beta distribution we will mean the transformed version defined above.

From now on we will apply Equations 5 and 7, to determine the compression efficiency of SM and VLB methods for random matrices generated as describe above.

With $w = 0$, a single Beta distribution is used. In Figure 1, we show some distributions of bit-lengths for some combinations of α_1 and β_1 . From the figure it can be seen that a large variety of unimodal distributions can be generated in the interval $[1, 64]$.

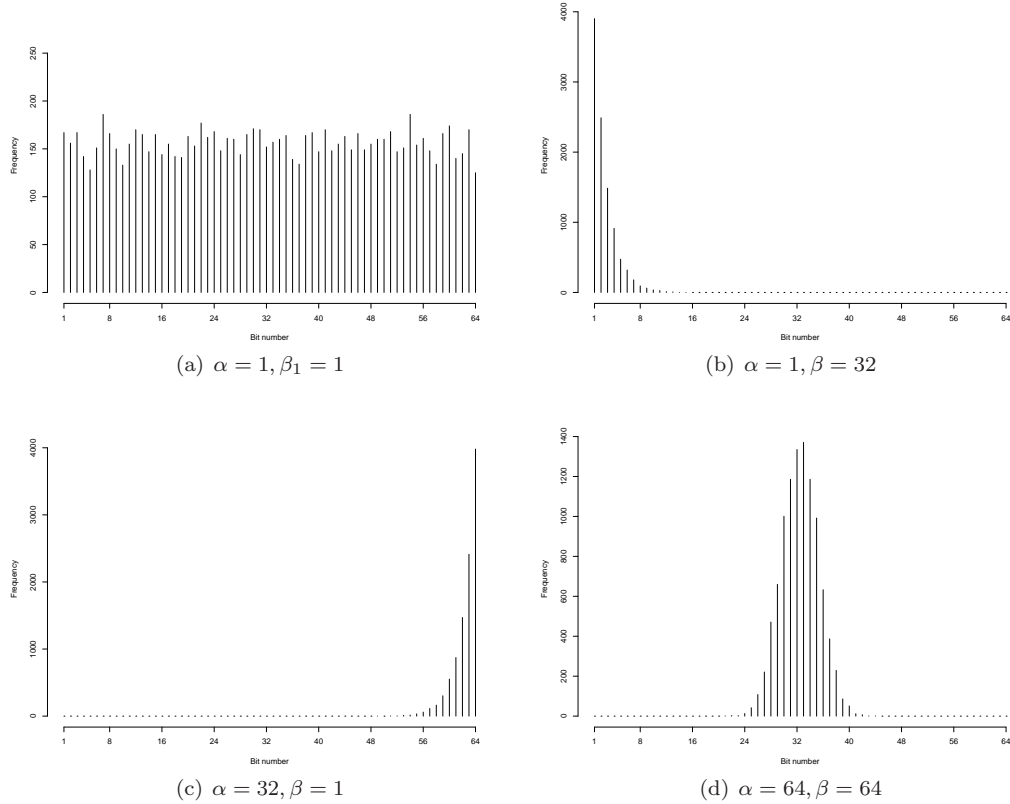


Figure 1. Histograms constructed from samples with 10,000 elements, generated from a Beta distribution. Below each histogram is possible to verify the parameters used.

As we are sampling from a large set distributions of bit-length, represented by the mixture of betas presented above, in order to make our results more general, we will base our analysis on the expected bit-length of a sample, since the efficiency of both methods depends on it. So, from Equations 5 and 7, the expected efficiencies become:

$$E(\eta_1) = 1 - \frac{k}{64} \quad (10)$$

$$E(\eta_2) = 1 - \frac{E(b)}{64} - \frac{k}{64} \quad (11)$$

where k in 11, is set to 7 (the bit-length required to represent the largest possible bit-length: 64). In 10, k is the bit-length of the greatest element, or in the worst case, 64.

We will use the difference $D = E(\eta_1) - E(\eta_2)$ to compare the efficiency of the two methods. Thus a positive D will favor SM method while a negative D favors VLB method.

The expected compression efficiency in the following numeric experiments, will be calculated from 3 matrices of dimension 10000, generated as described, and presented in tables and figures below.

In Figure 2, we can see the distribution of efficiencies and their difference for a sample generated from a single Beta distribution of bit-lengths. Note that both methods can achieve efficiencies greater than 80% for matrices with very small numbers. Also note that the VLB method is more efficient in the majority of cases.




Figure 2. Comparing compression efficiency of methods 1 and 2. Color scale in (a), (b) and (c) represent average bit-length. In (a) we can see the difference $D = \eta_1 - \eta_2$. It can be seen that for most combination of α and β , $D < 0$, meaning the second method is more efficient to compress a sample of numbers with bit-lengths coming from a $Beta(\alpha, \beta)$ distribution. However, there is a small region in parameter space, which is shown in white on (d), where the first method is more efficient. This region corresponds to the dots in red in (a), where the average bit-length is higher. In panels (b) and (c), we can see the efficiencies of SM and VLB methods, respectively.

Now let $w = 0.5$, i.e., matrices will have elements with bit-lengths coming from a mixture of beta

distributions, $B_1 \sim \text{Beta}(\alpha_1, \beta_1)$ and $B_2 \sim \text{Beta}(\alpha_2, \beta_2)$. The expected value for this mixture is shown in Equation 12.

$$E(B) = 0.5E(B_1) + 0.5E(B_2) \quad (12)$$

For bit-lengths coming from a mixture ($w > 0$), let the expected efficiencies for the SM and VLB methods be as given by Equations 13 and 14. So now, instead of having the efficiency be a function of greatest bit-length in the sample (denoted as k in 5 and 7), it will be a function of $\max\{E(B_1), E(B_2)\}$.

$$E(\eta_1) = 1 - \frac{\max\{E(B_1), E(B_2)\}}{64} \quad (13)$$

$$E(\eta_2) = 1 - 0.5\frac{E(B_1)}{64} - 0.5\frac{E(B_2)}{64} - \frac{\max\{E(B_1), E(B_2)\}}{64} \quad (14)$$

As before, we generate 3 matrices of dimension 10,000 for each parameterization, calculate the average efficiencies (Equations 13 and 14) and their difference D .

Before moving on to efficiency results and analyses, let's first inspect samples from the mixture of transformed Beta distributions. Figures 3 and 4, show a few parameterizations and their resulting sample distributions. It is important to note that from the mixture we can now generate bimodal distributions as well as the unimodal types tested before. Since we are making statements about efficiency as a function of the expected bit-length, it is important to verify if these statements hold for bimodal distributions as well.

After sampling uniformly ($[1, 5, 9, \dots, 64]$, $n = 65536$) the bit-length space and comparing efficiencies, we summarized the results on Table 3. In it we see how many parameterizations (from our sample) favor each method. We can also look at the distribution of efficiencies on our samples for each method (figure 5), which clearly demonstrate the greater expected efficiency of method VLB (figure 5(b)).

Table 2. Efficiency comparison of SM and VLB methods for parameters covering uniformly the support of B . Column n shows the number of parameter combinations with which each method has superior compression.

Methods	n	Method wins
SM	592	0.9034%
VLB	64944	99.0966%
Total	65536	100%

As we have shown, the VLB method is more effective compressing most integer datasets up to 64 bits in size. This is due to its ability to exploit the variance in the data set and reduce the waste of bits in the representation of some numbers. In specific cases where the variance in the data null or too small, method I will be more efficient. As a matter of fact, for matrices where all elements have the same bit-length, SM method will always be better, regardless of bit-length (figures 6(a) and 6(b)), The only exception if for bit-length 64 where neither method is able to compress the data.

Discussion

Calculating Efficiencies

To determine the best compression method to apply, it's necessary to inspect the distribution of bit-lengths of matrix elements. When matrix elements are small or have nearly-constant bit-length, the SM Method is best, otherwise, the VLB method should be chosen.

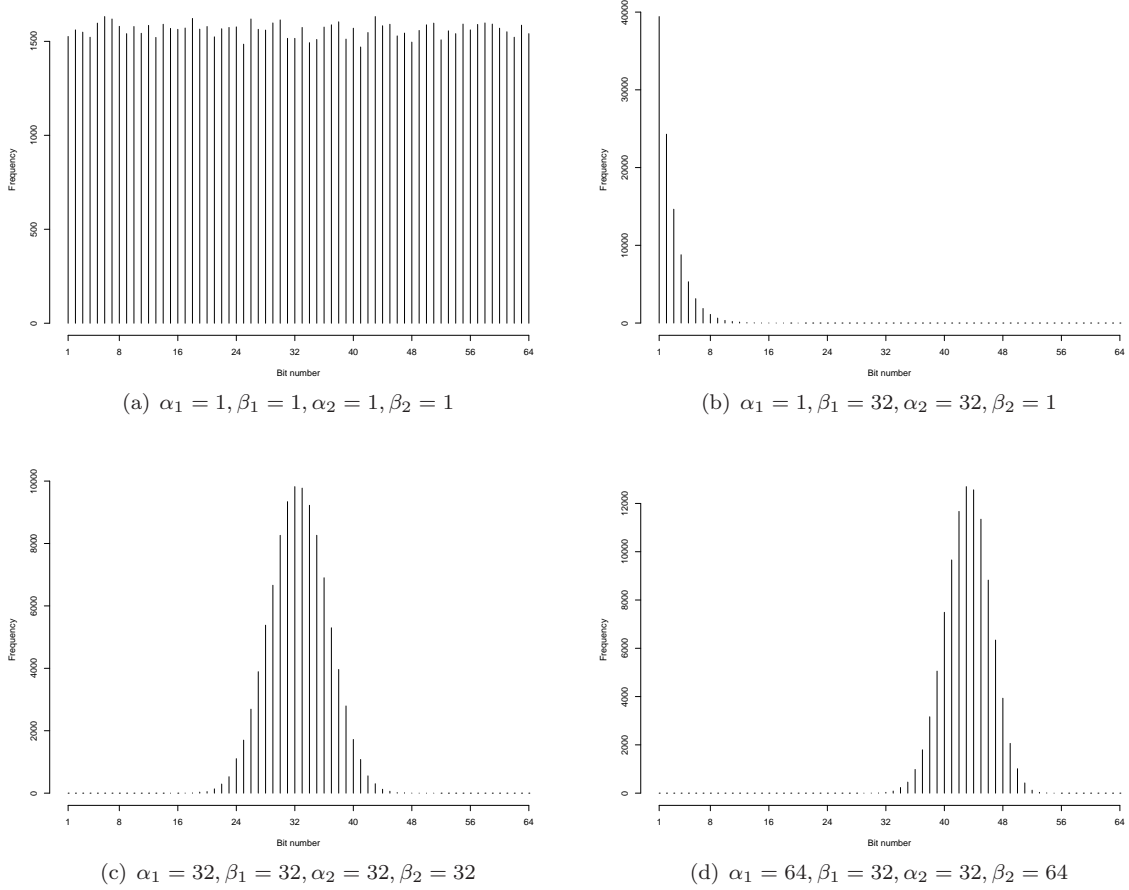


Figure 3. Histograms constructed from samples with 10,000 elements, generated from the combination of two distributions Betas with $\lambda = 0.5$. Below each histogram is possible to verify the parameters used..

As an example, let $M_{r \times c}$ be a integer matrix such that the half of its elements have bit-length 1 and the other half 64. Recalling Equation 7, now we have two groups of elements (by bit-length), $b_1 = 1$, $b_2 = 64$ and $f_i = \frac{rc}{2}$ for $i = 1$ and 2. As the greatest bit-length is 64, then $k = 7$. Compression efficiency η_2 can be calculated using Equation 7. After plugging in our numbers, we obtain a compression of 38.29%.

$$\eta_2 = 1 - \frac{\sum_{i=1}^2 b_i \times f_i}{64 \times rc} - \frac{7}{64}$$

$$\eta_2 = 1 - \frac{1 \times \frac{rc}{2} + 64 \times \frac{rc}{2}}{64 \times rc} - \frac{7}{64}$$

$$\eta_2 = 1 - \frac{32.5 \times rc}{64 \times rc} - \frac{7}{64}$$

$$\eta_2 \approx 1 - 0.5078 - 0.1093$$

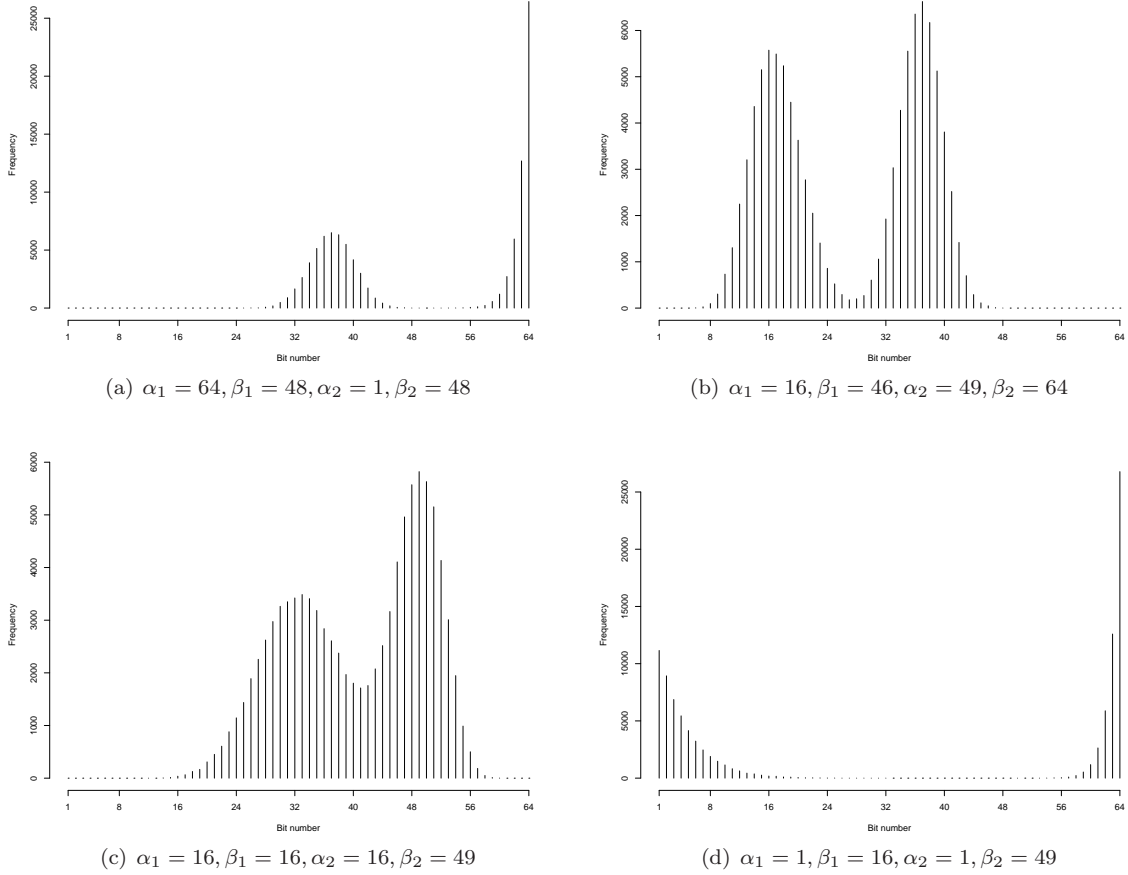


Figure 4. Histograms constructed from samples with 10,000 elements, generated from the combination of two distributions Betas with $\lambda = 0.5$. Below each histogram is possible to verify the parameters used.

$$\eta_2 \approx 38.29\%$$

The efficiency of the VLB method is influenced by the relative size of the bit-length groups. In this first example we considered only two groups, each comprised of half the matrix elements. Let's now vary the relative frequency of the groups, $\frac{f_i}{rc}$, while sticking to two groups. Let's also assume that $\frac{f_i}{rc}$ is a good approximation to the probability of a given bit-length in a matrix, which we will denote by p_i .

With this definition we can rewrite the Equation 7, which becomes 15. In Equation 15, the $\frac{f_i}{rc}$ is replaced by p_i , representing the probability of elements from group i in matrix M .

$$\eta_2 = 1 - \frac{\sum_{i=1}^g b_i \times p_i}{64} - \frac{k}{64} \quad (15)$$

with

$$p_i = \frac{f_i}{rc} \quad (16)$$

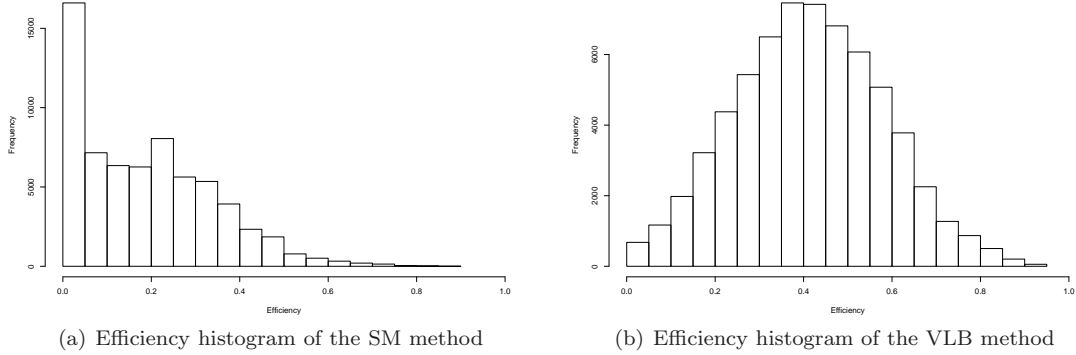


Figure 5. Efficiency histograms of the SM and VLB methods. Note that the VLB method has a greater average efficiency than SM method .

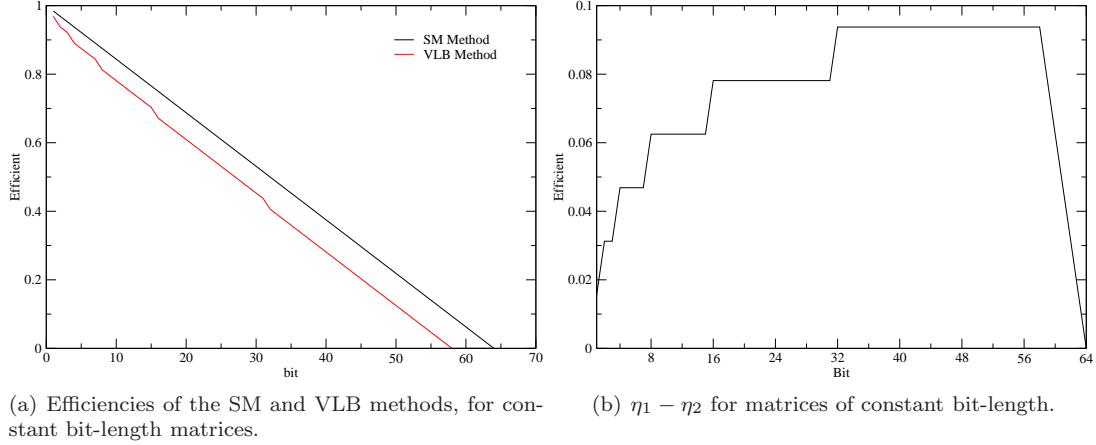


Figure 6. Compression efficiency of the SM and VLB methods for matrices of constant bit-length.

With the Equation 15 can analyze the influence of bit-length probability in compression efficiency. In this example, p_1 and p_2 represent the probability of elements of bit-lengths 1 and 64, respectively. Thus, efficiency is defined in Equation 17.

$$\eta_2 = 1 - \frac{1 \times p_1 + 64 \times p_2}{64} - \frac{7}{64} \quad (17)$$

Now, we can determine which probabilities give us the best and worst compression levels. When $\eta_2 = 1$, then the efficiency is maximal and if $\eta_2 = 0$, a efficiency is minimal. To calculate the values of p_1 and p_2 for both extreme values of η_2 , we must solve the linear systems shown in Equations 18 and 19. The first equation on both systems come from the law of total probability. The second comes from 17 after setting η_2 to 1 and 0, respectively.

$$\eta_2 = 1 : \begin{cases} p_1 + p_2 = 1 \\ p_1 + 64p_2 = 7 \end{cases} \quad (18)$$

Solving the system above, we find that when $p_1 = 0.9047$ and $p_2 = 0.0953$, efficiency is maximal, and

in this particular case is equal to 87.5%.

$$\eta_2 = 0 : \begin{cases} p_1 + p_2 = 1 \\ p_1 + 64p_2 = 57 \end{cases} \quad (19)$$

Thus, when $p_1 = 0.1111$ and $p_2 = 0.8889$ the efficiency is minimal for the VLB method. For other combinations see Table 3. Looking at this table, one can see two negative efficiencies, when (p_1, p_2) assume the values (0,1) and (0.1,0.9). This correspond th cases when the method increases the memory requirements instead of decreasing it.

Table 3. Combinations p_1 and p_2 to calculate the efficiency.

p_1	p_2	η_2
0.0	1.0	-0.109
0.1	0.9	-0.010
0.2	0.8	0.087
0.3	0.7	0.185
0.4	0.6	0.284
0.5	0.5	0.382
0.6	0.4	0.481
0.7	0.3	0.579
0.8	0.2	0.678
0.9	0.1	0.776
1.0	0.0	0.875

So far, we have examined only two groups (hence two probabilities) of bit-length for the sake of simplicity. Before we generalize to probability distributions let's take a quick look at the efficiencies for more groups, with uniform probability:

- 3 groups with bit-lengths 1, 32 and 64 bits, efficiency $\eta_2 = 0.3854$,
- 5 groups with bit-lengths 1, 16, 32, 48 and 64 bits, efficiency $\eta_2 = 0.3875$,
- 8 groups with bit-lengths 1, 8, 16, 24, 32, 40, 48, 56 and 64 bits, efficiency $\eta_2 = 0.3888$

When the distribution of the group probabilities is uniform, i.e., the groups have approximately the same size, efficiency is basically the same, regardless of the number of groups.

Now we can leverage the notion of bit-length probabilities, and study efficiency when bit-lengths follow some commonly used discrete probability distributions: Discrete Uniform, Binomial and Poisson. For all the experiments, we assume $k = 7$, that is, the maximum possible bit-length is 64 bits. Thus, efficiency obtained will not be the best possible, since for that we would need assume small values of k (Equation 15).

Discrete Uniform

Let the bit-lengths of the matrices be distribute according to the Uniform distribution $U(a = 1, b = 64)$, which means bit-lengths may take values in the set $\{1, 2, 3, \dots, 64\}$ with equal probability, i.e., $\frac{1}{64}$.

Theoretical Efficiency: Let the random variable $B \sim U(a = 1, b = 64)$ represent the bit-length of the elements of matrix M . Then $E(b_i) = \sum_i b_i \times p(b_i) = \frac{a+b}{2}$. Applying this result to the expected compression efficiency of VLB method (Equation 15), we have

$$E(\eta_2) = 1 - \frac{E(B)}{64} - \frac{k}{64} \quad (20)$$

assuming all bit-lengths are possible, i.e., $a = 1$ and $b = 64$, and hence $k = 7$, we can calculate η_2 :

$$E(\eta_2) = 1 - \frac{\frac{1+64}{2}}{64} - \frac{7}{64} \approx 38.28\% \quad (21)$$

This result agrees with the numerical estimates presented in Table 4.

Numerical Estimates: To calculate the VLB efficiency, we generated a matrices with 100 ($M_{10 \times 10}$), 10,000 ($M_{100 \times 100}$) and 1,000,000 ($M_{1,000 \times 1,000}$) elements with 1, 8, 16, 32 and 64 bit numbers. The average efficiency (Table 4) is calculated from a 1,000 replicates of each matrix size. possvel verificar que quanto menor for o nmero de bits necessrios para representar os elementos da matriz, maior ser o nivel de compactao.

Table 4. Compression efficiency of VLB method of samples with bit-lengths coming from a Discrete Uniform distribution $U(a = 1, b = 64)$. Average efficiency ($\bar{\eta}_2 \pm$ standard deviation) were calculated over a 1,000 replicates.

Expected bit-length	Matrix sizes		
	Sample size		
	100	10,000	1,000,000
1	0.8750±0.0000	0.8750±0.0000	0.8750±0.0000
8	0.8202±0.0031	0.8203±0.0003	0.8203±0.0000
16	0.7580±0.0066	0.7578±0.0007	0.7578±0.0001
32	0.6330±0.0142	0.6329±0.0014	0.6328±0.0001
64	0.3826±0.0288	0.3828±0.0028	0.3828±0.0003

Binomial Distribution

For the binomial distribution, we will use $Bin(n, p)$, with the number of trials n representing the greatest possible bit-length in the matrix, with np giving us an expected bit-length.

Theoretical Efficiency: Let bit-length (B) be a random variable with Binomial distribution, $B \sim B(n = 64, p = 0.5)$, $E(b_i) = \sum_i b_i \times p(b_i) = n \times p$ and the efficiency becomes (with $k = 7$):

$$\begin{aligned} E(\eta_2) &= 1 - \frac{64 \times p}{64} - \frac{k}{64} \\ &= 1 - \frac{64 \times 0.5}{64} - \frac{7}{64} \approx 39.05\% \end{aligned} \quad (22)$$

Which again agrees with estimates in Table 5.

Numerical Estimates: For these experiments, the parameter n represents the maximum bit-length of matrix elements and takes values in $\{1, 8, 16, 32, 64\}$. In this case, we evaluate the efficiency as a function of the parameter n , and matrix size. Even though efficiency does not depend on matrix size, we tried different sizes to test the stability of the compression algorithm. Results are shown in Table 5. As expected, smaller bit-lengths lead to higher compression efficiencies.

Table 5. Compression efficiency with bit-lengths distributed according to a binomial distribution $B(n, 0.5)$. Parameter $n \in \{1, 8, 16, 32, 64\}$ represents the maximum bit-length. Since $p = 0.5$ the expected bit-length is $n/2$ (first column)

Expected bit-length ($n/2$)	Efficiency		
	Sample size		
	100	1,000	1,000,000
1	0.8828±0.0004	0.8828±0.0004	0.8828±0.0000
8	0.8283±0.0022	0.8281±0.0002	0.8281±0.0000
16	0.7656±0.0032	0.7656±0.0003	0.7656±0.0000
32	0.6406±0.0045	0.6406±0.0004	0.6406±0.0000
64	0.3910±0.0065	0.3906±0.0006	0.3906±0.0001

Poisson Distribution

With bit-length derived from a $Poisson(\lambda)$, the parameter λ corresponds to the expected bit-length. For the purpose of this analysis this Poisson distribution is truncated at 64.

Theoretical Efficiency: Let bit-length $B \sim Poisson(\lambda = 32)$. In this case, $E(b_i) = \lambda$, with $k = 7$, the efficiency becomes:

$$E(\eta_2) = 1 - \frac{\lambda}{64} - \frac{k}{64} \quad (23)$$

$$= 1 - \frac{32}{64} - \frac{7}{64} \approx 39.06\% \quad (24)$$

This result is in accordance to Table 6.

Numerical Estimates: The results for this simulation can be seen in 6. Note that increasing the bit number to represent numbers increases, there is a loss of efficiency in the compression process. Another interesting fact is the emergence of a negative efficiency. In this case, for values generated with a $\lambda = 64$, the second method requires more memory to allocate the matrix formed by elements generated from this distribution.

Table 6. Compression efficiency with bit-lengths distributed according to a Poisson distribution (λ), where λ represents the expected bit-length.

Expected bit-length (λ)	Efficiency η_2		
	Matrix Size		
	100	1,000	1,000,000
1	0.8751±0.0015	0.8750±0.0004	0.8750±0.0000
8	0.7654±0.0046	0.7656±0.0005	0.7656±0.0000
16	0.6405±0.0064	0.6406±0.0006	0.6406±0.0001
32	0.3908±0.009	0.3906±0.0009	0.3906±0.0001
64	0.4086±0.0456	0.4073±0.0046	0.4072±0.0005

These results show that a good compression is guaranteed when bit-lengths are distributed according to the tested distributions regardless of sample size.

1 Conclusion

In this paper, we have focused in the compression of matrix data, since this is one of the most important application the authors foresee. However, the compression methodology presented can be applied to any numerical data structure, with gains to performance and memory footprint[citar tese Crysttian e possveis artigos derivados].

Further discussions about doing computation with such compressed data-structures will be the subject of another manuscript (in preparation) in which we will present details about the implementation of the compression algorithm, and benchmarks on classical linear algebra tasks such as those in Linpack[ref].

Representation of floating point numbers is also possible within the proposed compression framework, but at the expense of precision in their representation. Although this may sound like a limitation, when we take into consideration that most experimental data have fewer “significant” digits than the maximal precision available in modern computers, fairly good compression may still be achievable for floats.

Acknowledgments

We would like to thank Claudia Torres Codeço, Paulo Cezar Carvalho and Moacyr Silva for fruitful discussions and key ideas which helped improve the manuscript.

References

1. Salomon D, Motta G, Bryant D (2009) Handbook of Data Compression. London: Springer, 1359 pp.
2. Salomon D (2007) Data Compression: The Complete Reference. Number v. 10 in Data compression: the complete reference. London: Springer-Verlag New York Incorporated.
3. Lynch C (2008) Big data: How do your data grow? Nature 455: 28–29.
4. Gaber M, Zaslavsky A, Krishnaswamy S (2005) Mining data streams: a review. ACM Sigmod Record 34: 18–26.
5. Yemliha T, Chen G, Ozturk O, Kandemir M, Degalahal V (2007) Compiler-directed code restructuring for operating with compressed arrays. In: Proceedings of the 20th International Conference on VLSI Design held jointly with 6th International Conference: Embedded Systems. Citeseer, pp. 221–226.
6. Dodson D, Grimes R, Lewis J (1991) Sparse extensions to the fortran basic linear algebra subprograms. ACM Transactions on Mathematical Software (TOMS) 17: 253–263.
7. Van Der Walt S, Colbert S, Varoquaux G (2011) The numpy array: a structure for efficient numerical computation. Computing in Science & Engineering 13: 22–30.
8. Kane MJ, Emerson JW (2012) bigmemory: Manage massive matrices with shared memory and memory-mapped files. URL <http://CRAN.R-project.org/package=bigmemory>. R package version 4.3.0.
9. Bodden E, Clasen M, Kneis J (2007) Arithmetic coding revealed. In: Sable Technical Report 2007-5, Sable Research Group, School of Computer Science, (McGill University, Montréal).

Figure Legends

Tables